

PLAN-FM Bridge @ AAAI 2026 | 21 Jan 2026

Teaching LLMs to Plan: From Chain-of-Thought Instruction Tuning to Collaborative Constraint Translation

Pulkit Verma



pulkitverma.net | pulkitv@cse.iitm.ac.in



How do we combine human expertise with machine speed and scale, especially when lives are on the line?

The Research Gap

- Automated planners are powerful but require expert knowledge.
- Humans have intuition, but can't solve complex problems fast enough.
- LLMs are unreliable for critical decision-making.

What if we could have a system where domain experts could guide AI planning in natural language and trust that the results are valid?

Can we make LLMs reliable planners
AND
use them to make planning accessible?

PDDL-Instruct: Enhancing Symbolic Planning Capabilities in LLMs through Logical Chain-of-Thought Instruction Tuning



Pulkit Verma^{*}



Ngoc La^{*}



Anthony Favier^{*}



Swaroop Mishra[†]



Julie A. Shah^{*}



ICAPS 2025 Workshop on Planning in the Era of LLMs (LM4Plan @ ICAPS25)

LLMs are good at Reasoning

Article

DeepSeek-R1 incentivizes reasoning in LLMs through reinforcement learning

<https://doi.org/10.1038/s41586-025-09422-z>
Received: 14 February 2025
Accepted: 17 July 2025
Published online: 17 September 2025
Open access
Check for updates

General reasoning represents a long-standing and formidable challenge in artificial intelligence (AI). Recent breakthroughs, exemplified by large language models (LLMs)^{1,2} and chain-of-thought (CoT) prompting³, have achieved considerable success on foundational reasoning tasks. However, this success is heavily contingent on extensive human-annotated demonstrations and the capabilities of models are still insufficient for more complex problems. Here we show that the reasoning abilities of LLMs can be incentivized through pure reinforcement learning (RL), obviating the need for human-labelled reasoning trajectories. The proposed RL framework facilitates the emergent development of advanced reasoning patterns, such as self-reflection, verification and dynamic strategy adaptation. Consequently, the trained model achieves superior performance on verifiable tasks such as mathematics, coding competitions and STEM fields, surpassing its counterparts trained through conventional supervised learning on human demonstrations. Moreover, the emergent reasoning patterns exhibited by these large-scale models can be systematically used to guide and enhance the reasoning capabilities of smaller models.

Reasoning capability, the cornerstone of human intelligence, enables complex cognitive tasks ranging from mathematical problem-solving to logical deduction and programming. Recent advances in AI have demonstrated that LLMs can exhibit emergent behaviours, including reasoning abilities, when scaled to a sufficient size^{1,2}. However, achieving such capabilities in pre-training typically demands substantial

whereas unrestricted RL training can better incentivize the development of new reasoning capabilities in LLMs. Through this process, in the next section, our model (referred to as DeepSeek R1) naturally solves diverse and sophisticated reasoning problems, while solving reasoning problems, the model exhibits a tendency towards longer responses, incorporating verification, reflection and

Chain-of-Thought Prompting Elicits Reasoning in Large Language Models

Jason Wei Xuezhi Wang Dale Schuurmans Maarten Bosma
Brian Ichter Fei Xia Ed H. Chi Quoc V. Le Denny Zhou
Google Research, Brain Team
{jasonwei, dennyzhou}@google.com

Abstract

We explore how generating a *chain of thought*—a series of intermediate reasoning steps—significantly improves the ability of large language models to perform complex reasoning. In particular, we show how such reasoning abilities emerge naturally in sufficiently large language models via a simple method called *chain-of-thought prompting*, where a few chain of thought demonstrations are provided as exemplars in prompting. Experiments on three large language models show that chain-of-thought prompting improves performance on a range of arithmetic, commonsense, and symbolic reasoning tasks. The empirical gains can be striking. For instance, prompting a PaLM 540B with just eight chain-of-thought exemplars achieves state-of-the-art accuracy on the GSM8K benchmark of math word problems, surpassing even finetuned GPT-3 with a verifier.

Standard Prompting

Chain-of-Thought Prompting

nature medicine



Article <https://doi.org/10.1038/s41591-024-03423-7>

Toward expert-level medical question answering with large language models

Received: 14 June 2024
Accepted: 14 November 2024
Published online: 8 January 2025
Check for updates

A list of authors and their affiliations appears at the end of the paper

Large language models (LLMs) have shown promise in medical question answering, with Med-PaLM being the first to exceed a ‘passing’ score in United States Medical Licensing Examination style questions. However, challenges remain in how to form medical question answering and handling

edges these domain bounding M2 scores up / over 19%, dMCQA, human M2 answers ed-PaLM2 isor across s designed -world to generalist re still ‘aLM2 to tential in

ensing Examination

Successive Prompting for Decomposing Complex Questions

Dheeru Dua¹ Shivanshu Gupta¹ Sameer Singh^{1,2} Matt Gardner¹
¹University of California, Irvine, USA ²Allen Institute for Artificial Intelligence
¹Microsoft Semantic Machines
{ddua, shivag5, sameer}@uci.edu, mattgardner@microsoft.com

Abstract

Answering complex questions that require making latent decisions is a challenging task, especially when limited supervision is available. Recent works leverage the capabilities of large language models (LLMs) to perform complex question answering in a few-shot setting by

Who kicked the longest field goal in the first half?

Q: What are all the field goals in the first half?

A: 12-yard, 42-yard and 33-yard

Q: What is the largest value in: 15, 20, 45, 30, and 10?

ReTA: Recursively Thinking Ahead to Improve the Strategic Reasoning of Large Language Models

Jinhao Duan¹ Shiqi Wang² James D. Hendler³ Lichao Sun⁴
Tianlong Chen^{5,6,7} Bhavya Kailkhura¹ Kaidi Xu¹
¹Drexel University ²AWS AI Lab
³Lawrence Livermore National Laboratory
⁴Lehigh University ⁵UNC Chapel Hill ⁶MIT ⁷Harvard University

Abstract

Current logical reasoning evaluations of Large Language Models (LLMs) primarily focus on single-turn and static environments, such as arithmetic problems. The crucial problem of multi-turn, strategic reasoning is under-explored. In this work, we analyze the multi-turn strategic reasoning of LLMs through text-driven complete- and incomplete-

evaluations still focus on the linguistic capabilities of LLMs, e.g., reading understanding, without much strategic thinking. Therefore, beneath the impressive linguistic capabilities of LLMs, a critical question that has piqued the curiosity of researchers and practitioners alike: “*what lies beyond static logical reasoning for LLMs?*” Strategic multi-turn reasoning tasks, such as board and card games, test more reflection of real

LLMs are bad at Planning

Position: LLMs Can’t Plan, But Can Help Planning in LLM-Modulo Frameworks

Subbarao Kambhampati¹ Karthik Valmeekam¹ Lin Guan¹ Murthy Venkatakrishnan¹ Koushik Suresh¹
Siddhant Bhambril¹ Lucas Saldyt¹ Anil Mehta¹

Abstract

We argue that auto-regressive LLMs cannot, by themselves, do planning or self-verification (which is after all a form of reasoning), and shed some light on the reasons for misunderstandings in the literature. We also argue that LLMs should be viewed as universal approximate knowledge sources that have much more meaningful roles

with System 2 can seem to ring true, i are best seen as a g (see Figure 1). Ev a system that takes cannot possibly be Not surprisingly, li formance of LLMs

Large Language Models Still Can’t Plan (A Benchmark for LLMs on Planning and Reasoning about Change)

Alberto Olmo^{*}
School of Computing & AI
Arizona State University, Tempe,
aolmo@asu.edu

Subbarao Kambhampati
School of Computing & AI
Arizona State University, Tempe,
rao@asu.edu

On the Planning Abilities of Large Language Models : A Critical Investigation

Karthik Valmeekam
School of Computing & AI
Arizona State University, Tempe,
kvalmeek@asu.edu

Sarath Sreedharan^{*}
Department of Computer Science,
Colorado State University, Fort Collins,
sarath.sreedharan@colostate.edu

Abstract

Intrigued by the claims of emergent reasoning general web corpora, in this paper, we set abilities. We aim to evaluate (1) the effect autonomously in commonsense planning task source of heuristic guidance for other agent We conduct a systematic study by generating lar to the ones employed in the International LLMs in two distinct modes: *autonomous* : LLMs’ ability to generate executable plans a best model (GPT-4) having an average success However, the results in the heuristic mode mode, we demonstrate that LLM-generated for underlying sound planners and additional help provide feedback on the generated plan plan generation.

Can Large Language Models Really Improve by Self-critiquing Their Own Plans?

Karthik Valmeekam^{*}
School of Computing & AI
Arizona State University, Tempe,
kvalmeek@asu.edu

Matthew Marquez^{*}
School of Computing & AI
Arizona State University, Tempe,
mmarqu22@asu.edu

Subbarao Kambhampati
School of Computing & AI
Arizona State University, Tempe,
rao@asu.edu

Abstract

There have been widespread claims about Large Language Models (LLMs) being able to successfully verify or self-critique their candidate solutions in reasoning problems in an iterative mode. Intrigued by those claims, in this paper we set out to investigate the verification/self-critiquing abilities of large language models in the context of planning. We evaluate a planning system that employs LLMs for both plan generation and verification. We assess the verifier LLM’s performance against ground-truth verification, the impact of self-critiquing on plan generation, and the influence of varying feedback levels on system performance. Using GPT-4

Can we leverage LLMs' reasoning capabilities for Planning?

- What works for reasoning in LLMs?
- How to leverage it for planning?

Planning Domain Definition Language (PDDL)

```
(:action pickup
```


```
  :parameters (?ob)
```

```
  :precondition (and  
    (handempty)  
    (ontable ?ob))
```



Precondition: This condition must be true for this action to execute

```
  :effect (and  
    (not (handempty))  
    (not (ontable ?ob))  
    (holding ?ob))
```



Effect: This is a set of conditions, one of which becomes true when this action is executed


```
)
```

What works for reasoning in LLMs?

- Finetuning
- Instruction tuning (finetuning with instructions)
- Chain-of-Thought prompting


Finetuning

Adapt a pre-trained general LLM to excel at a specific task (planning) by training on domain examples.




Finetuning with Negative Examples

Add some failing plans, label them as incorrect, and add them to the finetuning data.




Instruction Finetuning

Add Explanations: Instructions teach the model WHAT planning means, not just PATTERNS in data. Tell it to check preconditions, apply effects, and verify goals.




Augment finetuned LLM with Chain-of-Thought Prompting



Making the model show intermediate reasoning steps for planning instead of jumping to the final answer.

PDDLInstruct



Reasoning Chain Optimization

optimize the model parameters θ_t to
improve the generation of high-quality reasoning chains

$$\theta_t^r = \theta_t - \delta_1 \nabla_{\theta_t} L_{\text{reasoning}}(\theta_t, \mathbb{D}_{\text{reasoning}}^t)$$

$\{(s_{i-1}, a_i, s_i, f_i) : \forall \text{ steps in CoT plans generated at iteration } t\}$

loss function that measures the quality of the generated reasoning chains

Reasoning Chain Optimization: $\theta_t^r = \theta_t - \delta_1 \nabla_{\theta_t} L_{reasoning}(\cdot)$

This objective encourages the model to produce step-by-step reasoning that correctly:

1. checks all necessary preconditions before applying actions;
2. tracks state changes resulting from action effects; and
3. detects logical inconsistencies in proposed plans.

Reasoning Chain Optimization

$$\mathcal{L}_{reasoning}(\theta_t, \mathbb{D}_{reasoning}^t) =$$

$$\frac{1}{|\mathbb{D}_{reasoning}^t|} \sum_{(s_{i-1}, a_i, s_i, f_i) \in \mathbb{D}_{reasoning}^t} d(s_i, s_i^{expected}) + \lambda_{feedback} \mathcal{L}_{feedback}$$

$$\mathcal{L}_{feedback} = \begin{cases} 0 & \text{if action } a_i \text{ is valid} \\ \alpha_{pre} & \text{if precondition violation detected} \\ \alpha_{eff} & \text{if incorrect effect applied} \\ \alpha_{goal} & \text{if goal not achieved} \end{cases}$$

End-Task (Final) Performance Optimization

optimize from the reasoning-improved parameters θ_t^r to enhance overall planning

$$\theta_{t+1} = \theta_t^r - \delta_2 \nabla_{\theta_t^r} \mathcal{L}_{final}(\theta_t^r, \mathbb{D}_{final}^t)$$

$\{(d_j, p_j, \pi_i^t, v_j^t) : \forall \text{ problems } j \text{ at iteration } t\}$

loss function that measures how well the final outputs match the expected answers in the training data

End-Task Performance Optimization: $\theta_{t+1} = \theta_t^r - \delta_2 \nabla_{\theta_t^r} L_{final}(\cdot)$

This objective ensures that improvements in logical reasoning translate to practical planning capability of producing accurate plans.

Empirical Evaluation: Objectives

RQ1: Does logical CoT instruction tuning improve plan validity compared to standard approaches?

RQ2: How does the quality of feedback (binary vs. detailed) affect planning performance?

RQ3: How well does the approach generalize across different planning domains?

Empirical Evaluation: Dataset and Models

Three Domains:

- Blockworld
- Logistics
- Mystery Blocksworld

Three Models:

- Llama-3-8B
- GPT-4
- Gemma-3-270M

Benchmark

PlanBench: An Extensible Benchmark for Evaluating Large Language Models on Planning and Reasoning about Change

Karthik Valmeekam
School of Computing & AI
Arizona State University, Tempe.
kvalmeek@asu.edu

Matthew Marquez
School of Computing & AI
Arizona State University, Tempe.
mmarqu22@asu.edu

Alberto Olmo
School of Computing & AI
Arizona State University, Tempe.
aolmoher@asu.edu

Sarath Sreedharan*
Department of Computer Science,
Colorado State University, Fort Collins.
sarath.sreedharan@colostate.edu

Subbarao Kambhampati
School of Computing & AI
Arizona State University, Tempe.
rao@asu.edu

Logical CoT instruction tuning improves Plan Validity

Model	Domain	Baseline	Only P1	Only P2	PDDL-INSTRUCT				
				Detailed	Binary		Detailed		
				$\eta = 15$	$\eta = 10$	$\eta = 15$	$\eta = 10$	$\eta = 15$	
Llama-3	Blocksworld	28%	78%	72%	84%	89%	91%	94%	
	Mystery BW	1%	32%	17%	47%	49%	59%	64%	
	Logistics	11%	23%	45%	61%	72%	75%	79%	
GPT-4	Blocksworld	35%	41%	76%	79%	84%	87%	91%	
	Mystery BW	3%	17%	19%	39%	44%	54%	59%	
	Logistics	6%	27%	51%	64%	69%	72%	78%	
Gemma-3	Blocksworld	7%	12%	19%	37%	39%	54%	56%	
	Mystery BW	0%	2%	3%	22%	28%	24%	28%	
	Logistics	2%	13%	11%	18%	33%	27%	43%	

Detailed feedback is better than Binary Feedback

Model	Domain	Baseline	Only P1	Only P2	PDDL-INSTRUCT			
				Detailed	Binary		Detailed	
				$\eta = 15$	$\eta = 10$	$\eta = 15$	$\eta = 10$	$\eta = 15$
Llama-3	Blocksworld	28%	78%	72%	84%	89%	91%	94%
	Mystery BW	1%	32%	17%	47%	49%	59%	64%
	Logistics	11%	23%	45%	61%	72%	75%	79%
GPT-4	Blocksworld	35%	41%	76%	79%	84%	87%	91%
	Mystery BW	3%	17%	19%	39%	44%	54%	59%
	Logistics	6%	27%	51%	64%	69%	72%	78%
Gemma-3	Blocksworld	7%	12%	19%	37%	39%	54%	56%
	Mystery BW	0%	2%	3%	22%	28%	24%	28%
	Logistics	2%	13%	11%	18%	33%	27%	43%

PDDLInstruct's improved performance generalizes across domains

Model	Domain	Baseline	Only P1	Only P2	PDDL-INSTRUCT			
				Detailed	Binary		Detailed	
				$\eta = 15$	$\eta = 10$	$\eta = 15$	$\eta = 10$	$\eta = 15$
Llama-3	Blocksworld	28%	78%	72%	84%	89%	91%	94%
	Mystery BW	1%	32%	17%	47%	49%	59%	64%
	Logistics	11%	23%	45%	61%	72%	75%	79%
GPT-4	Blocksworld	35%	41%	76%	79%	84%	87%	91%
	Mystery BW	3%	17%	19%	39%	44%	54%	59%
	Logistics	6%	27%	51%	64%	69%	72%	78%
Gemma-3	Blocksworld	7%	12%	19%	37%	39%	54%	56%
	Mystery BW	0%	2%	3%	22%	28%	24%	28%
	Logistics	2%	13%	11%	18%	33%	27%	43%

Limitations

- Optimizing instruction tuning data.
- Fine-grained analysis of planning performance.
- Comparison with SoTA symbolic planners.
- Extending domain coverage.



A Collaborative Numeric Task Planning Framework based on Constraint Translations using LLMs



Anthony Favier



Ngoc La




Pulkit Verma



Julie A. Shah



ICAPS 2025 Workshop on Planning in the Era of LLMs (LM4Plan @ ICAPS25)



Formal Automated Planning

Requires:

- programming knowledge.
- technical expert interventions.




Improve Planning Accessibility

- Avoid “intuitively bad solutions” and focus
- Explore specific strategies in a “Let’s try this and rollback” approach.
- Dynamically refine solutions and iterate toward more effective outcomes.




Hybrid Collaborative Planning Framework



A **symbolic** automated planner computes plans


An **LLM-based system** acts as an **interface** and for model elicitation

Hybrid Collaborative Planning Framework



Human can **influence** problem solving,
without technical expertise requirements


Main capabilities: Chat, Suggestions, Translation



Chat Capability

Chat


- Get insight on the problem
- Summarize problem
- Modify existing plans
- **No PDDL** for user



Providing Suggestions Capability

**Highlight information,
Make suggestions**


- Retrieve **external** information (RAG) and real-time APIs.
- Can generate real-time **weather** constraints, not modeled in original PDDL problem



Translation Capability

**Main contribution:
Planning + Translation**

- **Translate** human guidance into planning constraints
- The updated problem is **solved** by **symbolic** planner
- **Simulator** to **visualize** plan



No constraints

Setting: DEFAULT
Planning mode: anytime, TO=15.0
Problem (zenoreal):

- NumericTCORE/benchmark/ZenoTravel-no-constraint/domain_with_n.pddl
- PDDL/zenoreal.pddl

=== ADDING CONSTRAINT ===

Enter your constraint:

1) Human input

Elapsed Time: 0.0 s

Confirm

Translate

Risk Analysis

Chat

Plan

Plans

Previous:

None

Current:

None

R0 - Only use plane1

- D1- Plane2 cannot board any passengers
- D2- Plane2 cannot debark any passengers
- D3- Plane2 cannot fly slow between any cities
- D4- Plane2 cannot fly fast between any cities
- D5- Plane2 cannot refuel
- D6- Plane3 cannot board any passengers
- D7- Plane3 cannot debark any passengers
- D8- Plane3 cannot fly slow between any cities
- D9- Plane3 cannot fly fast between any cities
- D10- Plane3 cannot refuel

2) Added Constraints

- Plane3 cannot fly slow between any cities
- Plane3 cannot fly fast between any cities
- Plane3 cannot refuel

Are you satisfied with the decomposition? If not, provide any desired feedback or type 'explain'.

User: yes

Encoding ...

Elapsed Time: 0.0 s

Confirm

Translate

Risk Analysis

Chat

Plan

Plans

Previous:

None

Current:

None

R0 - Only use plane1

- D1- Plane2 cannot board any passengers
- D2- Plane2 cannot debark any passengers
- D3- Plane2 cannot fly slow between any cities
- D4- Plane2 cannot fly fast between any cities
- D5- Plane2 cannot refuel
- D6- Plane3 cannot board any passengers
- D7- Plane3 cannot debark any passengers
- D8- Plane3 cannot fly slow between any cities
- D9- Plane3 cannot fly fast between any cities
- D10- Plane3 cannot refuel

3) Plan

- PDDL/zenoreal.pddl

Constraints loaded

=== PLANNING ===

Compiling ... OK [1.52s]

Planning (anytime, TO=15.0s) ... OK [15.06s]

Plans

Previous:

None

Current:

Plan-Length: 48
Metric: 15536.0
Planning time: 15.06
Found Plan:
0.0: (refuel_plane1)
1.0: (board_person4_plane1_boston)
2.0: (flyfast_plane1_boston_washington)
3.0: (board_person2_plane1_washington)
4.0: (board_person8_plane1_washington)
5.0: (flyslow_plane1_washington_boston)
6.0: (refuel_plane1)
7.0: (flyslow_plane2_washington_washington)
8.0: (flyslow_plane1_boston_dallas)
9.0: (board_person9_plane1_dallas)
10.0: (flyfast_plane1_dallas_seattle)
11.0: (debark_person9_plane1_seattle)
12.0: (refuel_plane1)
13.0: (flyslow_plane1_seattle_denver)
14.0: (debark_person4_plane1_denver)
15.0: (flyslow_plane1_denver_washington)
16.0: (refuel_plane1)
17.0: (flyslow_plane1_washington_seattle)
18.0: (flyslow_plane1_seattle_dallas)

Elapsed Time: 16.5 s

Confirm

Translate

Risk Analysis

Chat

Plan



ZenoR

PDSim



4) Simulation

Plan Actions

refuel (plane1)

flyslow (plane1, boston, denver)

flyfast (plane1, denver, washington)

refuel (plane1)


board (person8, plane1, washington)

flyslow (plane1, washington, denver)

debark (person8, plane1, denver)

flyslow (plane1, denver, boston)

board (person4, plane1, boston)



flyslow

located(plane1, denver)

Plan Panel

Action Tab

Speed Controls

Object Info Panel

Camera Controls

Elapsed Time: 16.5 s

Confirm

Translate

Risk Analysis


Chat

Plan

Guidance Translation

Translate user inputs as *guidance* for the *solver*

Two-step process:



Evaluation of translation quality: Ablation Study

4 Settings to evaluate our translation pipeline:

ECODING: LLM alone

+ **VERIFIER:** Symbolic syntax checker

+ **DECOMP:** Constraint decomposition

+ **HUMAN:** Human interventions on decomposition

Evaluation of translation quality: Ablation Study

Model:
Claude Sonnet 4
(thinking enabled)

Setting	Translation		Human interventions
	Parsable	Correct	Time (s)
Encoding	26	19	29.3 ± 12.3
+ Verifier	30	20	35.8 ± 13.5
+ Decomposition			
+ Human			

Table 1: Ablation study reporting syntax and semantic accuracy ($N = 30$)

Correct Syntax

- LLM alone makes syntax mistakes
- Symbolic verifier feedback fixes syntax mistakes

Evaluation of translation quality: Ablation Study

Model:
Claude Sonnet 4
(thinking enabled)

Setting	Translation			Human interventions
	Parsable	Correct	Time (s)	
Encoding	26	19	29.3 ± 12.3	0
+ Verifier	30	20	35.8 ± 13.5	0
+ Decomposition	30	20	55.0 ± 26.2	0
+ Human	30	27	81.9 ± 53.7	12

Table 1: Ablation study reporting syntax and semantic accuracy ($N = 30$)

Satisfying semantic accuracy

- Decomposition no direct effect
- But allows for human review
- Human intervention significantly improves correctness

Evaluation of translation quality: Ablation Study

Model:
Claude Sonnet 4
(thinking enabled)


Setting	Translation			Human interventions
	Parsable	Correct	Time (s)	
Encoding	26	19	29.3 ± 12.3	0
+ Verifier	30	20	35.8 ± 13.5	0
+ Decomposition	30	20	55.0 ± 26.2	0
+ Human	30	27	81.9 ± 53.7	12

Table 1: Ablation study reporting syntax and semantic accuracy ($N = 30$)

Seems faster than human experts

- Ours ~ 82 s (SD=53.7) vs. Prior work 180s (SD=78)
- But comparison maybe unfair
 - similar but not identical constraints

Hybrid collaborative planning framework (neuro-symbolic)



Creates a **collaborative, mixed-initiative planning** scheme where the human can influence problem solving, **without** technical expertise requirements

PLAN-FM Bridge @ AAAI 2026 | 21 Jan 2026

Teaching LLMs to Plan: From Chain-of-Thought Instruction Tuning to Collaborative Constraint Translation

Pulkit Verma

pulkitverma.net | pulkitv@cse.iitm.ac.in

